

# CECS 326-01

## Operating Systems

Dylan Dang (ID 026158052)

### Assignment 4

Due Date: 11/3/2022

Submission Date: 11/3/2022

### Program Description

Like the previous assignment, I am tasked make use of the POSIX implementation of the Linux shared memory mechanism. This program will incorporate the `fork()`, `exec()`, and `wait()` system calls to control the program that a child process is to execute. Since the program will make use of the POSIX implementation it will utilize the `shm_open()`, `ftruncate()`, `mmap()`, `munmap()`, `close()` functions. All these system calls/functions will be integrated into two C files: `master.c` and `slave.c`. Program will also utilize struct from given `myShm.h`. However, it won't stop there. This assignment tackles the potential problem for race condition due to the concurrent access and modification of the index variable in the shared data structure. Hence, mutual exclusion is necessary. I will be utilizing semaphore to enforce the necessary mutual exclusion. The named semaphore mechanism includes `sem_wait()`, `sem_post()`, `sem_open()`, `sem_close()`, and `sem_unlink()`.

#### Used function/system call & their definitions:

- `fork()` allows a new process to be created and consists of a copy of the address space of the original process. This mechanism allows the parent process to communicate easily with its child process.
- `exec()` replaces the process's memory space with a new program. The `exec()` system call loads a binary file into memory and starts its execution.
- `wait()` moves the parent process itself off the ready queue until the termination of the child.
- `shm_open()` returns a file descriptor that is associated with the shared "memory object" specified by name.
- `ftruncate()` truncates the file indicated by the open file descriptor to the indicated length.
- `mmap()` establishes a mapping between an address space of a process and a file associated with the file descriptor.

- `munmap()` remove any mappings for those entire pages containing any part of the address space of the process starting at 'address' (shm base address) and continuing for 'length' bytes.
- `close()` used to close a file, in this case it is used to close the shared memory segment as if it were a file.
- `sem_wait()` decrements by one the value of the semaphore. It will decrement when value is greater than zero. If value is zero, then the current thread will block until the value becomes greater than zero.
- `sem_post()` posts to a semaphore, incrementing its value by one. If the resulting value is greater than zero and if there is a thread waiting on the semaphore, the waiting thread decrements the value by one and continues running.
- `sem_open()` opens a named semaphore, returning a semaphore pointer that may be used on subsequent calls to `sem_post()`, `sem_wait()`, and `sem_close()`.
- `sem_close()` closes a named semaphore that was previously opened by a thread of the current process using `sem_open()`. This frees system resources associated with the semaphore on behalf of the process.
- `sem_unlink()` unlinks a name semaphore. The name of the semaphore is removed from the set of names used by the named semaphores.

## master.c description

Just like the previous assignment, Assignment 3, the *master.c* program is designed to make use of the `fork()`, `exec()`, and `wait()` system calls to create child processes with each child process executing a *slave* with its child number and the shared memory segment name passed to it from `exec()`. The given header file, `myShm.h` is also utilized within this program. The *master* program starts off with identifying itself and then storing the data from the commandline parameters such as the number of children to be created, shared memory segment name, and semaphore name. *Master* will use `shm_open()` to return a file descriptor which will be used to reference the content from the shared memory. `ftruncate()` is also used here to configure the size of the shared memory. After that, `mmap()` is used to map the memory region between an address space of a process and a file associated with the file descriptor.

Like in the case of assignment 3, this is the step where master uses `fork()` to create new child processes and `exec()` to pass arguments to the child; however, before that, there is a need to address the potential issue of race condition due to the concurrent access and modification of the index variable in the shared data structure. To combat this potential problem, I must make use of mutual exclusion to limit the number of processes operating on the index. Utilization of semaphores and a shared memory buffer is necessary to solve this dilemma. Master uses `sem_open()` to open a named semaphore which allows for the use of the other named semaphore mechanisms. This will be followed by a `sem_unlink()` call to indicate that the named semaphore will be destroyed once all processes have ceased. Now is the time for *master* to use `fork()` to create new child processes and `exec()` to pass arguments to the child. Each child process is to execute a *slave* with its child number and the shared memory

segment name passed to it from `exec()`. The *master* should output the number of slaves it has created and must wait for all of them to finish. `Sem_close()` will be used since the semaphore will not be used any longer in *master* and it also brings the benefit of freeing system resources. Upon termination of all child processes, *master* outputs the content of the shared memory segment filled in by the slaves. Finally, *master* uses the `munmap()` function to remove any mappings, close the shared memory segment using `close()`, removes the name of the shared memory segment using `shm_unlink()`, and then exits.

### slave.c description:

The *slave.c* program procedures are quite like *master.c* program. *Slave* program also utilizes the given header file, `myShm.h`. The program starts off with the *slave* identifying itself. It then proceeds to show its child number and the shared memory segment name it received from the `exec()` system call from *master*. Like the *master* program, the *slave* program also uses the `shm_open` function to return a file descriptor, the `ftruncate()` to configure the size of the shared memory, and the `mmap()` to create a map between an address space of a process and the file associated with the file descriptor.

Since semaphores are used in this assignment, *slave* will also use `sem_open()` to open a named semaphore using the name of the semaphore that was passed as an argument from master. *Slave* will then use `sem_wait()` to hold a semaphore or keep it waiting. This is beneficial since I am trying to prevent any race condition problems. The function checks the semaphore value and determines whether the semaphore should be blocked. This mechanism allows for *slave* to then writes its child number into the next available slot in the shared memory without the worry of another process accessing it at the same time. `Sem_post()` is used when the process is finished operating on the index since this function increments the semaphore value by 1 meaning the semaphore is unblocked and free. When all processes are done operating on the index, *slave* will use `sem_close()` to close the semaphore. Finally, *slave* uses the `munmap()` function to remove any mappings and uses `close()` to close the shared memory segment and then terminates.

### master.c, slave.c, compilation screenshot:

```
dylandang@Dylan-Blade:~/Assignment 4$ gcc -o master master.c -lrt -lpthread
dylandang@Dylan-Blade:~/Assignment 4$ gcc -o slave slave.c -lrt -lpthread
dylandang@Dylan-Blade:~/Assignment 4$ ls
master  master.c  myShm.h  slave  slave.c
```

### running master executable screenshot:

```
dylandang@Dylan-Blade:~/Assignment 4$ ./master 3 my_shm_name my_sem_name
Master begins execution
Master created a shared memory segment named my_shm_name
Master created a semaphore named my_sem_name
Master initializes index in the shared struct to 0
Master created 3 child processes to execute slave
Master waits for all child processes to terminate
Slave 1 begins execution
I am child number 1, received shared memory name, my_shm_name, and semaphore name, my_sem_name
Slave 1 acquires access to shared memory segment, and structures it according to struct CLASS
Slave 1 copies index to a local variable i
Slave 1 writes its child number in response[0]
Slave 1 increments index
Slave 1 closes access to shared memory and semaphore and terminates
Slave 1 exits
Slave 3 begins execution
I am child number 3, received shared memory name, my_shm_name, and semaphore name, my_sem_name
Slave 3 acquires access to shared memory segment, and structures it according to struct CLASS
Slave 3 copies index to a local variable i
Slave 3 writes its child number in response[1]
Slave 3 increments index
Slave 3 closes access to shared memory and semaphore and terminates
Slave 3 exits
Slave 2 begins execution
I am child number 2, received shared memory name, my_shm_name, and semaphore name, my_sem_name
Slave 2 acquires access to shared memory segment, and structures it according to struct CLASS
Slave 2 copies index to a local variable i
Slave 2 writes its child number in response[2]
Slave 2 increments index
Slave 2 closes access to shared memory and semaphore and terminates
Slave 2 exits
Master received termination signals from all 3 child processes
Content of shared memory segment filled by child processes:
--- content of shared memory ---
1
3
2
```

## master.c code:

```
C master.c
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/shm.h>
7  #include <sys/stat.h>
8  #include <sys/mman.h>
9  #include <sys/types.h>
10 #include <sys/wait.h>
11 #include <errno.h>
12 #include <semaphore.h>
13 #include "myShm.h"
14
15 int main(int argc, char** argv) {
16
17     /*Shared memory file descriptor*/
18     int shm_fileDes;
19
20     /*Takes user input from commandline argument such as
21     number of children, shared memory name, semaphore name and stores it*/
22     int numChildren = atoi(argv[1]);
23     const char *segmentName = argv[2];
24     const char *semaphoreName = argv[3];
25
26     /*Declare required size of memory request*/
27     const int SIZE = sizeof(struct CLASS);
28
29     /*Create base address as pointer to struct CLASS*/
30     struct CLASS *shm_baseAdd;
31
32     /*When executed, master outputs a message to identify itself*/
33     printf("Master begins execution\n");
34
35     /*Output shared memory segment name*/
36     printf("Master created a shared memory segment named %s\n", segmentName);
37
38     /*Output semaphore name*/
39     printf("Master created a semaphore named %s\n", semaphoreName);
40
41     /*Output starting index of response*/
42     printf("Master initializes index in the shared struct to 0\n");
43
44     /*The shm_open() function returns a file descriptor that is associated with the shared "memory object" specified by name.*/
45     shm_fileDes = shm_open(segmentName, O_CREAT | O_RDWR, 0666);
46
47     /*Checking for failure*/
48     if (shm_fileDes == -1) {
49         printf("ERROR: Shared memory failed; shm_open() failed\n");
50         exit(1);
51     }
52
53     /*The ftruncate() function truncates the file indicated by the open file descriptor to the indicated length. */
54     ftruncate(shm_fileDes, SIZE);
55
56     /*The mmap() function establishes a mapping between an address space of a process and a file associated with the file descriptor
57     at the offset for length of bytes.*/
58     shm_baseAdd = (struct CLASS *) mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fileDes, 0);
59
60     /*Checking for failure*/
61     if (shm_baseAdd == MAP_FAILED) {
62         printf("ERROR: Map failed; mmap() failed\n");
63         exit(1);
64     }
65
66     /*Update/Track index*/
67     shm_baseAdd->index = 0;
68
69     /*The sem_open() function opens a named semaphore, returning a semaphore pointer that may be used on
70     subsequent calls to sem_post(), sem_wait(), and sem_close().
71
72     The parameters of sem_open() are:
73     1. name - a pointer to the null-terminated name of the semaphore to be opened
74     2. oflag - option flags
75     3. mode - permission flags
76     4. value - initial value of the named semaphore */
77     sem_t *mutex_sem = sem_open(semaphoreName, O_CREAT, 0666, 1);
78 }
```

```

79  /*Checking for failure*/
80  if (mutex_sem == SEM_FAILED) {
81      printf("ERROR: sem_open failed(): %s\n", strerror(errno));
82      exit(1);
83  }
84
85  /*The sem_unlink() function unlinks a named semaphore.
86  The name of the semaphore is removed from the set of names used by named semaphores.
87
88  The parameter of sem_unlink() is:
89  1. name - a pointer to the null-terminated name of the semaphore to be unlinked */
90  if (sem_unlink(semaphoreName) == -1) {
91      printf("ERROR: sem_unlink() failed: %s\n", strerror(errno));
92      exit(1);
93  }
94
95  /*Determine number of child process to create*/
96  for (int i = 0; i < numChildren; i++) {
97
98      /*fork() system call is used to create child process*/
99      pid_t childPID = fork();
100
101      /*buffer for second argument in execlp()*/
102      char childNum[5];
103
104      /*Convert integer to char in order to be passed as an argument*/
105      sprintf(childNum, "%d", i + 1);
106
107      /*execute child process via execlp() function call*/
108      if (childPID == 0){
109          execlp("./slave", segmentName, childNum, semaphoreName, NULL);
110      }
111  }
112
113  /*Output number of child process created by master to be executed by slave*/
114  printf("Master created %d child processes to execute slave\n", numChildren);
115
116  /*Waiting for all child processes to exit*/
117  printf("Master waits for all child processes to terminate\n");
118
119  /*Output acknowledgement message*/
120  while(wait(NULL) != -1);
121  printf("Master received termination signals from all %d child processes\n", numChildren);
122
123  /*Output the content of the shared memory segment*/
124  printf("Content of shared memory segment filled by child processes:\n");
125  printf("--- content of shared memory ---\n");
126
127  for (int i = 0; i < numChildren; i++) {
128      /*For every child process, output the content of the shared memory filled in by the slave*/
129      printf("%d\n", shm_baseAddr -> response[i]);
130  }
131
132  /*The sem_close() function closes a named semaphore that was previously opened by a thread of the current process using sem_open().
133  The sem_close() function frees system resources associated with the semaphore on behalf of the process.sem_close()
134
135  The parameter for sem_close() is:
136  1. sem - a pointer to an opened named semaphore. This semaphore is closed for this process.*/
137  if (sem_close(mutex_sem) == -1) {
138      printf("ERROR from Master: sem_close failed: %s\n", strerror(errno));
139      exit(1);
140  }
141
142  /*The munmap() function shall remove any mappings for those entire pages containing
143  any part of the address space of the process starting at (shared memory base) address and continuing for (SIZE) length bytes*/
144  if (munmap(shm_baseAddr, SIZE) == -1) {
145      printf("ERROR from Master: Unmap failed; munmap() failed: %s\n", strerror(errno));
146      exit(1);
147  }
148
149  /*The close() function is used to close a file, in this case it is used to close the shared memory segment as if it was a file*/
150  if (close(shm_fileDes) == -1) {
151      printf("ERROR from Master: Close failed; close() failed: %s\n", strerror(errno));
152      exit(1);
153  }
154
155  /*Exit program*/
156  exit(0);
157

```

## slave.c code:

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/shm.h>
7  #include <sys/stat.h>
8  #include <sys/mman.h>
9  #include <sys/types.h>
10 #include <errno.h>
11 #include <semaphore.h>
12 #include "myShm.h"
13
14 int main(int argc, char **argv) {
15
16     /*Shared memory file descriptor*/
17     int shm_fileDes;
18
19     /*Takes user input from commandline argument such as
20     number of children, shared memory name, semaphore name and stores it*/
21     const char *segmentName = argv[0];
22     int childNum = atoi(argv[1]);
23     const char *semaphoreName = argv[2];
24
25     /*Declare required size of memory request*/
26     const int SIZE = sizeof(struct CLASS);
27
28     /*Local variable that holds next free index*/
29     int *i;
30
31     /*Create base address as pointer to struct CLASS*/
32     struct CLASS *shm_baseAdd;
33
34     /*Output identifying message*/
35     printf("Slave %d begins execution\n",childNum);
36
37     /*Output child number and segment name from execlp() system call */
38     printf("I am child number %d, received shared memory name, %s, and semaphore name, %s\n",childNum,segmentName,semaphoreName);
39
40     /*The shm_open() function returns a file descriptor that is associated with the shared "memory object" specified by name.*/
41     shm_fileDes = shm_open(segmentName, O_RDWR, 0666);
42
43     /*Checking for failure*/
44     if (shm_fileDes == -1) {
45         printf("ERROR from Slave %d. Shared Memory failed; shm_open() failed: %s\n",childNum, strerror(errno));
46         exit(1);
47     }
48
49     /*The mmap() function establishes a mapping between an address space of a process and a file associated with the file descriptor
50     at the offset for length of bytes.*/
51     shm_baseAdd = (struct CLASS *) mmap(0, SIZE,PROT_READ | PROT_WRITE, MAP_SHARED, shm_fileDes, 0);
52
53     /*Checking for failure*/
54     if (shm_baseAdd == MAP_FAILED) {
55         printf("ERROR from Slave %d. Map failed; shm_mmap() failed: %s\n",childNum, strerror(errno));
56         exit(1);
57     }
58
59     /*The sem_open() function opens a named semaphore, returning a semaphore pointer that may be used on
60     subsequent calls to sem_post(), sem_wait(), and sem_close().
61
62     The parameters of sem_open() are:
63     1. name - a pointer to the null-terminated name of the semaphore to be opened
64     2. oflag - option flags
65     3. mode - permission flags
66     4. value - initial value of the named semaphore */
67     sem_t *mutex_sem = sem_open(semaphoreName, O_CREAT, 0666, 1);
68
69     /*Checking for failure*/
70     if (mutex_sem == SEM_FAILED) {
71         printf("ERROR from Slave %d: sem_open failed(): %s\n",childNum, strerror(errno));
72         exit(1);
73     }
74
75     /*The sem_wait() function decrements by one the value of the semaphore.
76     The semaphore will be decremented when its value is greater than zero.
77     If the value of the semaphore is zero, then the current thread will block until the semaphore's value becomes greater than zero.
78
```



```

79 The parameter of sem_wait() is:
80 1. sem - a pointer to an initialized unnamed semaphore or opened named semaphore */
81 if (sem_wait(mutex_sem) == -1) {
82     printf("ERROR from Slave %d sem_wait() failed: %s\n",childNum, strerror(errno));
83     exit(1);
84 }
85
86 /*Output acknowledgement message*/
87 printf("Slave %d acquires access to shared memory segment, and structures it according to struct CLASS\n",childNum);
88
89 /*Update index value*/
90 i = &(shm_baseAdd->index);
91 printf("Slave %d copies index to a local variable i\n",childNum);
92
93 /*Write to shared memory*/
94 shm_baseAdd->response[i] = childNum;
95 printf("Slave %d writes its child number in response[%d]\n",childNum,*i);
96
97 /*Increment shared memory base address index*/
98 *i += 1;
99 printf("Slave %d increments index\n",childNum);
100
101 /*The sem_post() function posts to a semaphore, incrementing its value by one.
102 If the resulting value is greater than zero and if there is a thread waiting on the semaphore,
103 the waiting thread decrements the semaphore value by one and continues running.
104
105 The parameter of sem_post() is:
106 1. sem - a pointer to an initialized unnamed semaphore or opened named semaphore */
107 if (sem_post(mutex_sem) == -1) {
108     printf("ERROR from Slave %d: sem_post() failed: %s\n",childNum, strerror(errno));
109     exit(1);
110 }
111
112 /*The sem_close() function closes a named semaphore that was previously opened by a thread of the current process using sem_open().
113 The sem_close() function frees system resources associated with the semaphore on behalf of the process.sem_close()
114
115 The parameter for sem_close() is:
116 1. sem - a pointer to an opened named semaphore. This semaphore is closed for this process.*/
117 if (sem_close(mutex_sem) == -1) {
118     printf("ERROR from Slave %d: sem_close() failed: %s\n",childNum, strerror(errno));
119     exit(1);
120 }
121
122 /*The munmap() function shall remove any mappings for those entire pages containing
123 any part of the address space of the process starting at (shared memory base) address and continuing for (SIZE) length bytes*/
124 if (munmap(shm_baseAdd, SIZE) == -1) {
125     printf("ERROR from Slave %d. Unmap failed; munmap() failed; %s\n",childNum, strerror(errno));
126     exit(1);
127 }
128
129 /*The close() function is used to close a file, in this case it is used to close the shared memory segment as if it was a file*/
130 if (close(shm_fileDes) == -1) {
131     printf("ERROR from Slave %d. Close failed; close() failed: %s\n",childNum,strerror(errno));
132     exit(1);
133 }
134
135 printf("Slave %d closes access to shared memory and semaphore and terminates\n",childNum);
136 printf("Slave %d exits\n",childNum);
137
138 /*Exit program*/
139 exit(1);
140 }
141

```



myShm.h code:

```
h myShm.h
1  /* myShm.h*/
2  /* Header file to be used with master.c and slave.c
3  */
4  struct CLASS {
5
6      /*Index to next available response slot*/
7      int index;
8
9      /*Each child writes its child number here*/
10     int response[10];
11 };
12
```