

CECS 326-01

Operating Systems

Dylan Dang (ID 026158052)

Assignment 5

Due Date: 12/1/2022

Submission Date: 12/2/2022

Program Description

Like the previous assignment, I am tasked make use of the POSIX implementation of the Linux shared memory mechanism. This program will incorporate the `fork()`, `exec()`, and `wait()` system calls to control the program that a child process is to execute. Since the program will make use of the POSIX implementation it will utilize the `shm_open()`, `ftruncate()`, `mmap()`, `munmap()`, `close()` functions. All these system calls/functions will be integrated into two C files: `master.c` and `slave.c`. Program will also utilize `struct` from `myShm.h`. The last assignment tackled the potential problem for race condition due to the concurrent access and modification of the index variable in the shared data structure. However, for this assignment I am to utilize another semaphore to guard against race condition in the display device for output. Hence, mutual exclusion is necessary again. I will be utilizing semaphore to enforce the necessary mutual exclusion. The named semaphore mechanism includes `sem_wait()`, `sem_post()`, `sem_open()`, `sem_close()`, and `sem_unlink()`. The unnamed semaphore mechanism includes `sem_wait()`, `sem_post()`, `sem_init()`, and `sem_destroy()`.

Used function/system call & their definitions:

- `fork()` allows a new process to be created and consists of a copy of the address space of the original process. This mechanism allows the parent process to communicate easily with its child process.
- `exec()` replaces the process's memory space with a new program. The `exec()` system call loads a binary file into memory and starts its execution.
- `wait()` moves the parent process itself off the ready queue until the termination of the child.
- `shm_open()` returns a file descriptor that is associated with the shared "memory object" specified by name.
- `ftruncate()` truncates the file indicated by the open file descriptor to the indicated length.

- `mmap()` establishes a mapping between an address space of a process and a file associated with the file descriptor.
- `munmap()` remove any mappings for those entire pages containing any part of the address space of the process starting at 'address' (shm base address) and continuing for 'length' bytes.
- `close()` used to close a file, in this case it is used to close the shared memory segment as if it were a file.
- `sem_wait()` decrements by one the value of the semaphore. It will decrement when value is greater than zero. If value is zero, then the current thread will block until the value becomes greater than zero.
- `sem_post()` posts to a semaphore, incrementing its value by one. If the resulting value is greater than zero and if there is a thread waiting on the semaphore, the waiting thread decrements the value by one and continues running.
- `sem_open()` opens a named semaphore, returning a semaphore pointer that may be used on subsequent calls to `sem_post()`, `sem_wait()`, and `sem_close()`.
- `sem_close()` closes a named semaphore that was previously opened by a thread of the current process using `sem_open()`. This frees system resources associated with the semaphore on behalf of the process.
- `sem_unlink()` unlinks a name semaphore. The name of the semaphore is removed from the set of names used by the named semaphores.
- `sem_init()` initializes the unnamed semaphore at the address pointed to by `sem`.
- `sem_destroy()` destroys the unnamed semaphore at the address pointed to by `sem`.

master.c description

Just like the previous assignment, Assignment 4, the *master.c* program is designed to make use of the `fork()`, `exec()`, and `wait()` system calls to create child processes with each child process executing a *slave* with its child number and the shared memory segment name passed to it from `exec()`. The given header file, `myShm.h` is also utilized within this program. The *master* program starts off with identifying itself and then storing the data from the commandline parameters such as the number of children to be created, shared memory segment name, and semaphore name. *Master* will use `shm_open()` to return a file descriptor which will be used to reference the content from the shared memory. `ftruncate()` is also used here to configure the size of the shared memory. After that, `mmap()` is used to map the memory region between an address space of a process and a file associated with the file descriptor.

Like in the case of assignment 4, where mutual exclusion was used to stop multiple concurrent access to the index variable; we also use mutual exclusion to address the potential issue of race condition due to the concurrent output to the display device. Contrary to the previous assignment, I will make use of a named semaphore this time to combat the race condition regarding the access to the display device and I will use an unnamed semaphore and its mechanisms to combat race condition regarding access to the index variable. To begin, a named semaphore will be created to utilize `sem_wait()` and

`sem_post()`. These functions will reserve space for the processes' outputs to the display. This step is crucial since I want to make sure that the outputs are accurate to their respective processes. Next, I will be using `sem_init()` to initialize an unnamed semaphore. The unnamed semaphore is initialized at the address pointed to by `sem`, in this case is `mutex_sem` from our struct in `myShm.h`. This named semaphore will control the access to the index variable of the processes. Now is the time for *master* to use `fork()` to create new child processes and `exec()` to pass arguments to the child. Each child process is to execute a *slave* with its child number and the shared memory segment name passed to it from `exec()`. The *master* should output the number of slaves it has created and must wait for all of them to finish. `Sem_close()` will be used since the named semaphore will not be used any longer in *master* and it also brings the benefit of freeing system resources. `Sem_destroy()` will also be used to destroy the unnamed semaphore. Upon termination of all child processes, *master* outputs the content of the shared memory segment filled in by the slaves. Finally, *master* uses the `munmap()` function to remove any mappings, close the shared memory segment using `close()`, removes the name of the shared memory segment using `sem_unlink()`, and then `exits`.

slave.c description:

The *slave.c* program procedures are quite like *master.c* program. *Slave* program also utilizes the given header file, `myShm.h`. The program starts off with the *slave* identifying itself. It then proceeds to show its child number and the shared memory segment name it received from the `exec()` system call from *master*. Like the *master* program, the *slave* program also uses the `shm_open` function to return a file descriptor, the `ftruncate()` to configure the size of the shared memory, and the `mmap()` to create a map between an address space of a process and the file associated with the file descriptor.

Since both named semaphores and unnamed semaphores are used in this assignment, *slave* will also use `sem_open()` to open a named semaphore using the name of the semaphore that was passed as an argument from *master* and `sem_init()` will be used to create an unnamed semaphore from the struct. In the critical section, *slave* will then use `sem_wait()` to hold a semaphore or keep it waiting. This is beneficial since I am trying to prevent any race condition problems regarding the device display and the access to the index variable. The function checks the semaphore value and determines whether the semaphore should be blocked. This mechanism allows for *slave* to then writes its child number into the next available slot in the shared memory without the worry of another process accessing it at the same time. After leaving the critical section, `sem_post()` is used when the process is finished operating on the index and outputting to the display since this function increments the semaphore value by 1 meaning the semaphore is unblocked and free. When all processes are done operating on the index, *slave* will use `sem_close()` to close the named semaphore and `sem_destroy()` to destroy the unnamed semaphore. Finally, *slave* uses the `munmap()` function to remove any mappings and uses `close()` to close the shared memory segment and then terminates.